

```
In [ ]: import pandas as pd
import matplotlib.pyplot as plt
from sklearn.cluster import KMeans
from sklearn.preprocessing import MinMaxScaler, RobustScaler
from sklearn.metrics import silhouette_score
from sklearn.metrics import silhouette_samples
import matplotlib.cm as cm
import numpy as np
```

```
In [ ]: #reading in the dataset
df = pd.read_csv("../../data/healthcare-dataset-stroke-data.csv")
```

## Two-dimensional clustering

```
In [ ]: df = df.fillna(df.median())
```

```
In [ ]: df.head(3)
```

```
Out[ ]:
```

	id	gender	age	hypertension	heart_disease	ever_married	work_type	Residence_type	avg_glucose_level	bmi	smoking_status	stroke
0	9046	Male	67.0	0	1	Yes	Private	Urban	228.69	36.6	formerly smoked	1
1	51676	Female	61.0	0	0	Yes	Self-employed	Rural	202.21	28.1	never smoked	1
2	31112	Male	80.0	0	1	Yes	Private	Rural	105.92	32.5	never smoked	1

## BMI outlier handling

```
In [ ]: gluc_bmi = df[["avg_glucose_level", "bmi"]]
gluc_bmi.describe().round(2)
```

```
Out[ ]:
```

	avg_glucose_level	bmi
count	5110.00	5110.00
mean	106.15	28.86
std	45.28	7.70
min	55.12	10.30
25%	77.24	23.80
50%	91.88	28.10
75%	114.09	32.80
max	271.74	97.60

```
In [ ]: #Outlier handling
Q1 = df["bmi"].quantile(0.25)
Q3 = df["bmi"].quantile(0.75)
IQR = Q3 - Q1
gluc_bmi = gluc_bmi.loc[(gluc_bmi['bmi'] > (Q1-(1.5*IQR))) & (gluc_bmi['bmi'] < (Q3+(1.5*IQR)))]
```

```
In [ ]: gluc_bmi.describe().round(2)
```

```
Out[ ]:
```

	avg_glucose_level	bmi
count	4984.00	4984.00
mean	105.73	28.25
std	44.83	6.63
min	55.12	11.30
25%	77.20	23.60
50%	91.84	28.10
75%	113.76	32.30
max	271.74	46.20

```
In [ ]: #BMI tredshold
BMI_min = gluc_bmi["bmi"].describe()['min']
BMI_max = gluc_bmi["bmi"].describe()['max']
BMI_scaled = (25 - BMI_min)/(BMI_max-BMI_min)

#Avg_gluc_Level tredshold
gluc_min = gluc_bmi["avg_glucose_level"].describe()['min']
gluc_max = gluc_bmi["avg_glucose_level"].describe()['max']
gluc_scaled = (150 - gluc_min)/(gluc_max-gluc_min)
```

## Clusters by average glucose level and bmi

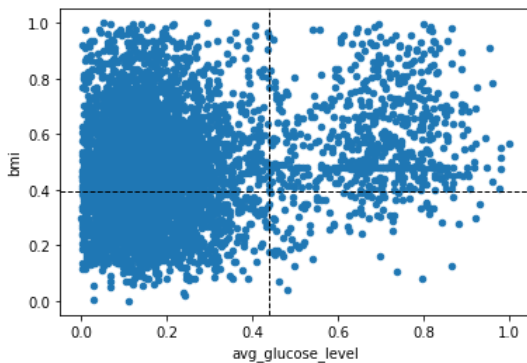
```
In [ ]: #robust scaling
#gluc_bmi = RobustScaler().fit_transform(gluc_bmi)
#gluc_bmi = pd.DataFrame(gluc_bmi, columns=["avg_glucose_level", "bmi"])
#gluc_bmi.describe().round(2)
```

```
In [ ]: #min-max scaling
#https://stats.stackexchange.com/questions/89809/is-it-important-to-scale-data-before-clustering
gluc_bmi = MinMaxScaler().fit_transform(gluc_bmi)
gluc_bmi = pd.DataFrame(gluc_bmi, columns=["avg_glucose_level", "bmi"])
gluc_bmi.describe().round(2)
```

```
Out[ ]:      avg_glucose_level      bmi
count      4984.00  4984.00
mean         0.23    0.49
std          0.21    0.19
min          0.00    0.00
25%          0.10    0.35
50%          0.17    0.48
75%          0.27    0.60
max          1.00    1.00
```

```
In [ ]: gluc_bmi.plot.scatter("avg_glucose_level", "bmi")
plt.axvline(gluc_scaled, color='k', linestyle='dashed', linewidth=1)
plt.axhline(BMI_scaled, color='k', linestyle='dashed', linewidth=1)
```

```
Out[ ]: <matplotlib.lines.Line2D at 0x23e86b1b0a0>
```



## Finding the right k in k-means

In order to find the best fitting number of clusters, we evaluate different k's by the elbow and silhouette method

### The elbow method


Distortion: Calculated as the avg of the squared distances from the cluster centers. Typically, the Euclidean distance metric is used.

```
In [ ]: distortions = []

K = range(2,15,1)

for k in K:
    kmeanModel = KMeans(n_clusters=k).fit(gluc_bmi)
    kmeanModel.fit(gluc_bmi)
    distortions.append(kmeanModel.inertia_)

plt.plot(K, distortions, 'bx-')
plt.xlabel('Values of K')
plt.ylabel('Distortion')
plt.title('The Elbow Method using Distortion')
plt.savefig("two_dim_elbow.png")
plt.show()
```



### The silhouette method

```
In [ ]: #https://scikit-learn.org/stable/auto_examples/cluster/plot_kmeans_silhouette_analysis.html

X = gluc_bmi

range_n_clusters = range(2,10,2)

for n_clusters in range_n_clusters:
    # Create a subplot with 1 row and 2 columns
```

```

fig, (ax1, ax2) = plt.subplots(1, 2)
fig.set_size_inches(12, 4)

# The 1st subplot is the silhouette plot
# The silhouette coefficient can range from -1, 1 but in this example all
# lie within [-0.1, 1]
ax1.set_xlim([-0.2, 1])

# The (n_clusters+1)*10 is for inserting blank space between silhouette
# plots of individual clusters, to demarcate them clearly.
ax1.set_ylim([0, len(X) + (n_clusters + 1) * 10])

# Initialize the clusterer with n_clusters value and a random generator
# seed of 10 for reproducibility.
clusterer = KMeans(n_clusters=n_clusters, random_state=10)
cluster_labels = clusterer.fit_predict(X)

# The silhouette_score gives the average value for all the samples.
# This gives a perspective into the density and separation of the formed
# clusters
silhouette_avg = silhouette_score(X, cluster_labels)
#print("For n_clusters =", n_clusters, "The average silhouette_score is :", silhouette_avg,)

# Compute the silhouette scores for each sample
sample_silhouette_values = silhouette_samples(X, cluster_labels)

y_lower = 10
for i in range(n_clusters):
    # Aggregate the silhouette scores for samples belonging to
    # cluster i, and sort them
    ith_cluster_silhouette_values = sample_silhouette_values[cluster_labels == i]

    ith_cluster_silhouette_values.sort()

    size_cluster_i = ith_cluster_silhouette_values.shape[0]
    y_upper = y_lower + size_cluster_i

    color = cm.nipy_spectral(float(i) / n_clusters)
    ax1.fill_betweenx(
        np.arange(y_lower, y_upper),
        0,
        ith_cluster_silhouette_values,
        facecolor=color,
        edgecolor=color,
        alpha=0.7,
    )

    # Label the silhouette plots with their cluster numbers at the middle
    ax1.text(-0.05, y_lower + 0.5 * size_cluster_i, str(i))

    # Compute the new y_lower for next plot
    y_lower = y_upper + 10 # 10 for the 0 samples

ax1.set_title("The silhouette plot for the various clusters.")
ax1.set_xlabel("The silhouette coefficient values")
ax1.set_ylabel("Cluster label")

# The vertical line for average silhouette score of all the values
ax1.axvline(x=silhouette_avg, color="red", linestyle="--")

ax1.set_yticks([]) # Clear the yaxis labels / ticks
ax1.set_xticks([-0.1, 0, 0.2, 0.4, 0.6, 0.8, 1])

# 2nd Plot showing the actual clusters formed
colors = cm.nipy_spectral(cluster_labels.astype(float) / n_clusters)
ax2.scatter(
    X["avg_glucose_level"], X["bmi"], marker=".", s=30, lw=0, alpha=0.7, c=colors, edgecolor="k"
)

# Labeling the clusters
centers = clusterer.cluster_centers_

# Draw white circles at cluster centers
ax2.scatter(
    centers[:, 0],
    centers[:, 1],
    marker="o",
    c="white",
    alpha=1,
    s=200,
    edgecolor="k",
)

for i, c in enumerate(centers):
    ax2.scatter(c[0], c[1], marker="%d$" % i, alpha=1, s=50, edgecolor="k")


ax2.set_title("The visualization of the clustered data.")
ax2.set_xlabel("Avg. glukose level")
ax2.set_ylabel("BMI")

plt.suptitle(
    "Silhouette analysis for KMeans clustering on sample data with n_clusters = %d"
    % n_clusters,
    fontsize=14,
    fontweight="bold",
)

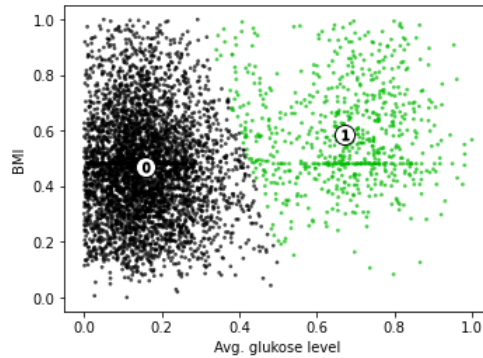
```

```
plt.savefig("two_dim_silhouette.png")
plt.show()
```


**Silhouette analysis for KMeans clustering on sample data with n\_clusters = 2**  
The silhouette plot for the various clusters.




The visualization of the clustered data.




**Silhouette analysis for KMeans clustering on sample data with n\_clusters = 4**  
The silhouette plot for the various clusters.




The visualization of the clustered data.



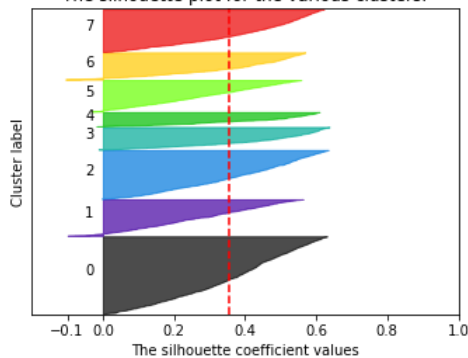
**Silhouette analysis for KMeans clustering on sample data with n\_clusters = 6**  
The silhouette plot for the various clusters.




The visualization of the clustered data.



**Silhouette analysis for KMeans clustering on sample data with n\_clusters = 8**  
The silhouette plot for the various clusters.



The visualization of the clustered data.



## Viewing the final k-means clusters


```
In [ ]: kmeans = KMeans(n_clusters=8)
kmeans.fit(gluc_bmi)
y_kmeans = kmeans.predict(gluc_bmi)

plt.scatter(gluc_bmi["avg_glucose_level"], gluc_bmi["bmi"], c=y_kmeans, s=10, cmap='viridis')
centers = kmeans.cluster_centers_
plt.scatter(centers[:, 0], centers[:, 1], c='white', s=200, alpha=0.9)

for i, c in enumerate(centers):
    plt.scatter(c[0], c[1], marker="%d$" % i, alpha=0.6, s=50, edgecolor="k")

plt.axvline(gluc_scaled, color='k', linestyle='dashed', linewidth=1)
plt.axhline(BMI_scaled, color='k', linestyle='dashed', linewidth=1)
```

```
Out[ ]: <matplotlib.lines.Line2D at 0x23e86ecf0a0>
```



## Cluster analysis

```
In [ ]: gluc_bmi_stroke = df[["avg_glucose_level", "bmi", "stroke"]]
```

```
In [ ]: #Outlier handling
Q1 = df["bmi"].quantile(0.25)
Q3 = df["bmi"].quantile(0.75)
IQR = Q3 - Q1
gluc_bmi_stroke = gluc_bmi_stroke.loc[(gluc_bmi_stroke['bmi'] > (Q1-(1.5*IQR))) & (gluc_bmi_stroke['bmi'] < (Q3+(1.5*IQR)))]
```

```
In [ ]: gluc_bmi_stroke = MinMaxScaler().fit_transform(gluc_bmi_stroke)
gluc_bmi_stroke = pd.DataFrame(gluc_bmi_stroke, columns=["avg_glucose_level", "bmi", "stroke"])
```

```
In [ ]: from matplotlib.lines import Line2D

kmeans = KMeans(n_clusters=8)
kmeans.fit(gluc_bmi_stroke[["avg_glucose_level", "bmi"]])
y_kmeans = kmeans.predict(gluc_bmi_stroke[["avg_glucose_level", "bmi"]])

fig, ax = plt.subplots(figsize=(14,8))

# add clusters to df
gluc_bmi_stroke["cluster"] = y_kmeans
# split df into strokes and non-strokes
non_strokes = gluc_bmi_stroke[gluc_bmi_stroke["stroke"] == 0]
strokes = gluc_bmi_stroke[gluc_bmi_stroke["stroke"] == 1]

# plot non-strokes
ax.scatter(non_strokes["avg_glucose_level"], non_strokes["bmi"], c=non_strokes["cluster"], s=30, marker='o', alpha=0.1, cmap='Spectral')

# plot strokes
scatter = ax.scatter(strokes["avg_glucose_level"], strokes["bmi"], c=strokes["cluster"], s=30, marker='D', cmap='Spectral')
centers = kmeans.cluster_centers_
ax.scatter(centers[:, 0], centers[:, 1], c='black', s=200, alpha=0.9)

# legend for cluster numbers
for i, c in enumerate(centers):
    plt.scatter(c[0], c[1], marker="$_d$_" % i, alpha=0.6, s=50, edgecolor="white")


# legend for stroke/no stroke
legend_elements = [Line2D([0], [0], marker='D', label='Stroke', markersize=10),
                   Line2D([0], [0], marker='o', label='No Stroke', markersize=10)]

legend2 = ax.legend(handles=legend_elements, loc='lower right', prop={'size':10})

plt.axvline(gluc_scaled, color='k', linestyle='dashed', linewidth=1)
plt.axhline(BMI_scaled, color='k', linestyle='dashed', linewidth=1)

plt.title(f"Cluster distribution with 8 clusters")
plt.savefig("two_dim_k8.png")
plt.show()
```

Cluster distribution with 8 clusters



```
In [ ]: # Calculate percentage of stroke cases for each cluster
total_strokes = strokes.shape[0]

for cluster in np.unique(y_kmeans):

    #share of data points
    shares = pd.DataFrame(gluc_bmi_stroke["cluster"].value_counts()/gluc_bmi_stroke.shape[0])
    shares = shares.sort_index()
    share = shares["cluster"][cluster]

    stroke_cases_cluster = strokes[strokes["cluster"] == cluster].shape[0]
    non_stroke_cases_cluster = non_strokes[non_strokes["cluster"] == cluster].shape[0]
    cluster_ratio = stroke_cases_cluster/(stroke_cases_cluster+non_stroke_cases_cluster)
    overall_stroke_percent = stroke_cases_cluster/total_strokes

    print(f"cluster {cluster}: {cluster_ratio:.2%} in cluster strokes, {overall_stroke_percent:.2%} of all strokes, {share:.2%} of all data po
```

cluster 0: 1.28% in cluster strokes, 3.66% of all strokes, 14.11% of all data points  
cluster 1: 2.90% in cluster strokes, 5.28% of all strokes, 9.01% of all data points  
cluster 2: 2.30% in cluster strokes, 4.88% of all strokes, 10.45% of all data points  
cluster 3: 5.79% in cluster strokes, 14.23% of all strokes, 12.14% of all data points  
cluster 4: 3.46% in cluster strokes, 11.38% of all strokes, 16.23% of all data points  
cluster 5: 5.05% in cluster strokes, 26.42% of all strokes, 25.82% of all data points  
cluster 6: 12.18% in cluster strokes, 11.79% of all strokes, 4.78% of all data points  
cluster 7: 14.78% in cluster strokes, 22.36% of all strokes, 7.46% of all data points

## Multidimensional clustering

```
In [ ]: data = df[["age", "avg_glucose_level", "heart_disease", "hypertension", "stroke"]]
scaler = MinMaxScaler()
data = scaler.fit_transform(data)
data = pd.DataFrame(data, columns=["age", "avg_glucose_level", "heart_disease", "hypertension", "stroke"])
data.head(3)
```


```
Out [ ]:    age  avg_glucose_level  heart_disease  hypertension  stroke
0  0.816895         0.801265             1.0             0.0       1.0
1  0.743652         0.679023             0.0             0.0       1.0
2  0.975586         0.234512             1.0             0.0       1.0
```

```
In [ ]: distortions = []

K = range(2,30,2)

for k in K:
    kmeanModel = KMeans(n_clusters=k).fit(data)
    kmeanModel.fit(data)
    distortions.append(kmeanModel.inertia_)

plt.plot(K, distortions, 'bx-')
plt.xlabel('Values of K')
plt.ylabel('Distortion')
plt.title('The Elbow Method using Distortion')
plt.savefig("mul_dim_elbow.png")
plt.show()
```



In [ ]: [https://scikit-learn.org/stable/auto\\_examples/cluster/plot\\_kmeans\\_silhouette\\_analysis.html](https://scikit-learn.org/stable/auto_examples/cluster/plot_kmeans_silhouette_analysis.html)

```
X = data

range_n_clusters = [5,10,15]

for n_clusters in range_n_clusters:
    fig, ax = plt.subplots()
    fig.set_size_inches(7, 4)
    ax.set_xlim([-0.2, 1])
    ax.set_ylim([0, len(X) + (n_clusters + 1) * 10])

    clusterer = KMeans(n_clusters=n_clusters, random_state=10)
    cluster_labels = clusterer.fit_predict(X)

    silhouette_avg = silhouette_score(X, cluster_labels)
    sample_silhouette_values = silhouette_samples(X, cluster_labels)

    y_lower = 10
    for i in range(n_clusters):

        ith_cluster_silhouette_values = sample_silhouette_values[cluster_labels == i]
        ith_cluster_silhouette_values.sort()

        size_cluster_i = ith_cluster_silhouette_values.shape[0]
        y_upper = y_lower + size_cluster_i

        color = cm.nipy_spectral(float(i) / n_clusters)
        ax.fill_betweenx(
            np.arange(y_lower, y_upper),
            0,
            ith_cluster_silhouette_values,
            facecolor=color,
            edgecolor=color,
            alpha=0.7,
        )

        # Label the silhouette plots with their cluster numbers at the middle
        ax.text(-0.05, y_lower + 0.5 * size_cluster_i, str(i))

        # Compute the new y_lower for next plot
        y_lower = y_upper + 10 # 10 for the 0 samples

    ax.set_title("The silhouette plot for the various clusters.")
    ax.set_xlabel("The silhouette coefficient values")
    ax.set_ylabel("Cluster label")

    # The vertical line for average silhouette score of all the values
    ax.axvline(x=silhouette_avg, color="red", linestyle="--")


    ax.set_yticks([]) # Clear the yaxis labels / ticks
    ax.set_xticks([-0.1, 0, 0.2, 0.4, 0.6, 0.8, 1])

    plt.suptitle(
        "Silhouette analysis for KMeans clustering (clusters = %d)"
        % n_clusters,
        fontsize=14,
        fontweight="bold",
    )

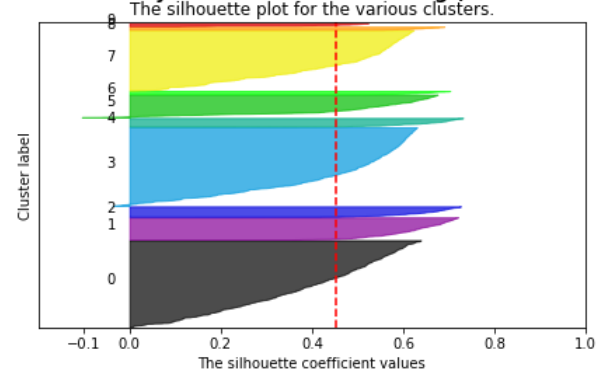
plt.savefig("mul_dim_silhouette.png")

plt.show()
```

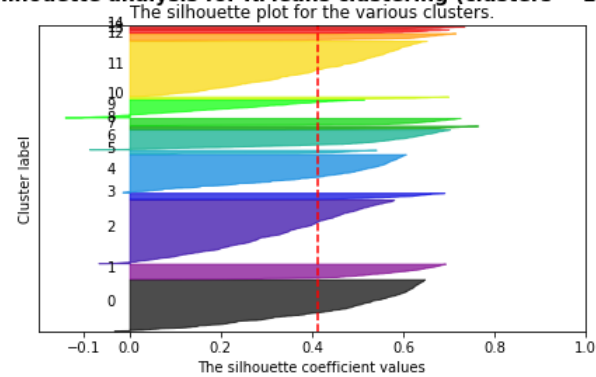
### Silhouette analysis for KMeans clustering (clusters = 5)



### Silhouette analysis for KMeans clustering (clusters = 10)



### Silhouette analysis for KMeans clustering (clusters = 15)



```
In [ ]: kmeans = KMeans(n_clusters=15)
kmeans.fit(data[["age", "avg_glucose_level", "heart_disease", "hypertension"]])
y_kmeans = kmeans.predict(data[["age", "avg_glucose_level", "heart_disease", "hypertension"]])
data["cluster"] = y_kmeans

# split into stroke and non stroke cases
non_strokes = data[data["stroke"] == 0]
strokes = data[data["stroke"] == 1]
```

```
In [ ]: def rev_min_max_func(scaled_val,target):
    max_val = max(df[target])
    min_val = min(df[target])
    val = (scaled_val*(max_val - min_val)) + min_val
    return val
```

```
In [ ]: data.head()
```

```
Out[ ]:
```

	age	avg_glucose_level	heart_disease	hypertension	stroke	cluster
0	0.816895	0.801265	1.0	0.0	1.0	5
1	0.743652	0.679023	0.0	0.0	1.0	3
2	0.975586	0.234512	1.0	0.0	1.0	11
3	0.597168	0.536008	0.0	0.0	1.0	12
4	0.963379	0.549349	0.0	1.0	1.0	7

```
In [ ]: for cluster in range(data["cluster"].max()+1):
    temp = data[data["cluster"]==cluster]
    print("cluster: " + str(cluster))
    print("age_mean: " + str(round(rev_min_max_func(temp["age"].mean(),"age"),2)))
    print("gluc_mean: " + str(round(rev_min_max_func(temp["avg_glucose_level"].mean(),"avg_glucose_level"),2)))
    print("heart_dis: " + str(temp["heart_disease"].sum()))
    print("hypertension: " + str(temp["hypertension"].sum()))
    print("size: " + str(temp.shape[0]) + ", strokes: " + str(temp["stroke"].sum()))
    print("in cluster share: " + str(round(((temp["stroke"].sum()/temp.shape[0])*100),2)) + "%")
    print("share of all strokes: " + str(round(((temp["stroke"].sum()/data["stroke"].sum()*100),2)) + "%")
    print("\n")
```

cluster: 0



age\_mean: 56.07  
gluc\_mean: 80.11  
heart\_dis: 0.0  
hypertension: 0.0  
size: 707, strokes: 24.0  
in cluster share: 3.39%  
share of all strokes: 9.64%

cluster: 1  
age\_mean: 23.28  
gluc\_mean: 80.3  
heart\_dis: 0.0  
hypertension: 0.0  
size: 664, strokes: 1.0  
in cluster share: 0.15%  
share of all strokes: 0.4%

cluster: 2  
age\_mean: 70.89  
gluc\_mean: 88.73  
heart\_dis: 0.0  
hypertension: 156.0  
size: 156, strokes: 24.0  
in cluster share: 15.38%  
share of all strokes: 9.64%

cluster: 3  
age\_mean: 65.78  
gluc\_mean: 213.81  
heart\_dis: 0.0  
hypertension: 0.0  
size: 277, strokes: 37.0  
in cluster share: 13.36%  
share of all strokes: 14.86%

cluster: 4  
age\_mean: 70.58  
gluc\_mean: 151.21  
heart\_dis: 64.0  
hypertension: 64.0  
size: 64, strokes: 13.0  
in cluster share: 20.31%  
share of all strokes: 5.22%

cluster: 5  
age\_mean: 68.74  
gluc\_mean: 211.75  
heart\_dis: 72.0  
hypertension: 0.0  
size: 72, strokes: 18.0  
in cluster share: 25.0%  
share of all strokes: 7.23%

cluster: 6  
age\_mean: 17.91  
gluc\_mean: 126.73  
heart\_dis: 0.0  
hypertension: 0.0  
size: 332, strokes: 0.0  
in cluster share: 0.0%  
share of all strokes: 0.0%

cluster: 7  
age\_mean: 64.85  
gluc\_mean: 204.89  
heart\_dis: 0.0  
hypertension: 142.0  
size: 142, strokes: 22.0  
in cluster share: 15.49%  
share of all strokes: 8.84%

cluster: 8  
age\_mean: 50.19  
gluc\_mean: 122.26  
heart\_dis: 0.0  
hypertension: 0.0  
size: 401, strokes: 14.0  
in cluster share: 3.49%  
share of all strokes: 5.62%

cluster: 9  
age\_mean: 6.43  
gluc\_mean: 85.83  
heart\_dis: 0.0  
hypertension: 0.0  
size: 581, strokes: 1.0  
in cluster share: 0.17%  
share of all strokes: 0.4%

cluster: 10  
age\_mean: 74.11  
gluc\_mean: 88.76  
heart\_dis: 0.0  
hypertension: 0.0  
size: 483, strokes: 60.0  
in cluster share: 12.42%

```
share of all strokes: 24.1%

cluster: 11
age_mean: 66.81
gluc_mean: 91.71
heart_dis: 140.0
hypertension: 0.0
size: 140, strokes: 16.0
in cluster share: 11.43%
share of all strokes: 6.43%
```

```
cluster: 12
age_mean: 36.62
gluc_mean: 198.48
heart_dis: 0.0
hypertension: 0.0
size: 138, strokes: 2.0
in cluster share: 1.45%
share of all strokes: 0.8%
```

```
cluster: 13
age_mean: 45.69
gluc_mean: 89.86
heart_dis: 0.0
hypertension: 136.0
size: 136, strokes: 7.0
in cluster share: 5.15%
share of all strokes: 2.81%
```

```
cluster: 14
age_mean: 38.7
gluc_mean: 84.56
heart_dis: 0.0
hypertension: 0.0
size: 817, strokes: 10.0
in cluster share: 1.22%
share of all strokes: 4.02%
```

## Visuals

```
In [ ]: #pip install plotly
```

```
In [ ]: import plotly
import plotly.graph_objs as go
plotly.offline.init_notebook_mode()
```

```
In [ ]: from sklearn.decomposition import PCA

# Create a PCA model to reduce our data to 2 dimensions for visualisation
pca = PCA(n_components=3)
pca.fit(data[["age", "avg_glucose_level", "heart_disease", "hypertension"]])

# Transform the scaled data to the new PCA space
X_reduced = pca.transform(data[["age", "avg_glucose_level", "heart_disease", "hypertension"]])

# Convert to df
reduced_df = pd.DataFrame(X_reduced, index=X.index, columns=['PC1', 'PC2', 'PC3'])

# # Add clusters and strokes
reduced_df["cluster"] = y_kmeans
reduced_df["stroke"] = data["stroke"]
reduced_df["heart_disease"] = data["heart_disease"]
reduced_df["hypertension"] = data["hypertension"]
```

```
In [ ]: reduced_df.head(3)
```

```
Out[ ]: 
```

	PC1	PC2	PC3	cluster	stroke	heart_disease	hypertension
0	0.508949	-0.710741	0.708314	5	1.0	1.0	0.0
1	0.163670	-0.236083	-0.088964	3	1.0	0.0	0.0
2	0.462517	-0.709081	0.517663	11	1.0	1.0	0.0


```
In [ ]: for x in range(0,max(reduced_df["cluster"])+1):
temp = reduced_df[reduced_df["cluster"]==x]
print("Cluster: " + str(x) + ", strokes: " + str(sum(temp["stroke"]))) + ", out of " + str(temp.stroke.count()))
```

```
Cluster: 0, strokes: 24.0, out of 707
Cluster: 1, strokes: 1.0, out of 664
Cluster: 2, strokes: 24.0, out of 156
Cluster: 3, strokes: 37.0, out of 277
Cluster: 4, strokes: 13.0, out of 64
Cluster: 5, strokes: 18.0, out of 72
Cluster: 6, strokes: 0.0, out of 332
Cluster: 7, strokes: 22.0, out of 142
Cluster: 8, strokes: 14.0, out of 401
Cluster: 9, strokes: 1.0, out of 581
Cluster: 10, strokes: 60.0, out of 483
Cluster: 11, strokes: 16.0, out of 140
Cluster: 12, strokes: 2.0, out of 138
Cluster: 13, strokes: 7.0, out of 136
Cluster: 14, strokes: 10.0, out of 817
```

```
In [ ]: # would also be nice to show the cluster centres,
#but we need to apply the same PCA transformation that we gave to the actual data points,
#so that they can appear in the new 2 dimensional PCA space
centres_reduced = pca.transform(kmeans.cluster_centers_)
```

```
In [ ]: import plotly.express as px
px.defaults.color_continuous_scale = px.colors.sequential.Rainbow
px.defaults.template = "ggplot2"

df = reduced_df
fig = px.scatter_3d(df, x='PC1', y='PC2', z='PC3', color='cluster')
fig.show()
```




## 2-d visual

```
In [ ]: # Plot the results (Only works for 2 dimensions)
fig, ax = plt.subplots(figsize=(12,8))

non_strokes = reduced_df[reduced_df["stroke"] == 0]
strokes = reduced_df[reduced_df["stroke"] == 1]

# plot non-strokes
ax.scatter(non_strokes["PC1"], non_strokes["PC2"], c=non_strokes["cluster"], s=30, marker='o', alpha=0.1, cmap='rainbow')
# and strokes
scatter = ax.scatter(strokes["PC1"], strokes["PC2"], c=strokes["cluster"], s=30, marker='D', cmap='rainbow')
# plot cluster centers
ax.scatter(centres_reduced[:, 0], centres_reduced[:, 1],
          marker='x', s=169, linewidths=3,
          color='r', zorder=10)
# legend for cluster numbers
legend1 = ax.legend(*scatter.legend_elements(),
                  loc="lower left", title="Clusters", prop={'size': 10})
ax.add_artist(legend1)
# legend for stroke/no stroke
legend_elements = [Line2D([0], [0], marker='D', label='Stroke', markersize=10),
                  Line2D([0], [0], marker='o', label='No Stroke', markersize=10),
                  ]
legend2 = ax.legend(handles=legend_elements, loc='upper left', prop={'size': 10})
plt.savefig("mul_dim_visu.png")
plt.show()
```




## PCA visuals



```
In [ ]: variance = pca.explained_variance_ratio_
fig, ax = plt.subplots()
ax.bar([0,1,2], variance)
ax.set_ylabel("Explained variance (%)")
ax.set_xlabel("Number of PCA")


print("explained variance summed: " + str(sum(variance).round(2)))
```

explained variance summed: 0.85



```
In [ ]: for c in range(3):
    abs(pca.components_)
    first_component = abs(pca.components_[c])
    x = ["age", "avg_gluc_lvl", "heart_disease", "hypertension"]
    fig, ax = plt.subplots(figsize=(6, 3))
    ax.set_ylim(0, 1)
    ax.bar(x, first_component)
    ax.set_ylabel("Importance of feature")
    fig.suptitle("Importance of features for PCA " + str(c))
```





In [ ]: